

02032021 - Advanced Pandas Functionalities

February 27, 2021

1 Pandas - Data Frame Operations

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: # create a simple dataframe
data = pd.DataFrame({
    'age' : [ 10, 22, 13, 21, 12, 11, 17],
    'section' : [ 'A', 'B', 'C', 'B', 'B', 'A', 'A'],
    'city' : [ 'Gurgaon', 'Delhi', 'Mumbai', 'Delhi', 'Mumbai', 'Delhi',
↳ 'Mumbai'],
    'gender' : [ 'M', 'F', 'F', 'M', 'M', 'M', 'F'],
    'favourite_color' : [ 'red', 'blue', 'yellow', 'blue', 'black', 'green',
↳ 'red']
})

# view the data
data
```

```
[2]:
```

	age	section	city	gender	favourite_color
0	10	A	Gurgaon	M	red
1	22	B	Delhi	F	blue
2	13	C	Mumbai	F	yellow
3	21	B	Delhi	M	blue
4	12	B	Mumbai	M	black
5	11	A	Delhi	M	green
6	17	A	Mumbai	F	red

1.1 loc and iloc

- **loc** gets rows (or columns) with particular **labels** from the index.
- **iloc** gets rows (or columns) at particular positions in the **index** (so it only takes integers).

1.1.1 Find all the rows based on any condition in a column

We can solve types of queries with a simple line of code using `pandas.DataFrame.loc[]`. We just need to pass the condition within the loc statement.

```
[3]: data.loc[data.age >= 15]
      # data[data.age >= 15]
```

```
[3]:   age section   city gender favourite_color
      1    22      B   Delhi      F           blue
      3    21      B   Delhi      M           blue
      6    17      A  Mumbai      F           red
```

1.1.2 Find all the rows with more than one condition

```
[4]: data.loc[(data.age >= 12) & (data.gender == 'M')]
      # data[(data.age >= 12) & (data.gender == 'M')]
```

```
[4]:   age section   city gender favourite_color
      3    21      B   Delhi      M           blue
      4    12      B  Mumbai      M          black
```

1.1.3 Select only required columns with a condition

```
[5]: data.loc[(data.age >= 12), ['city', 'gender', 'age']]
      # data[(data.age >= 12), ['city', 'gender', 'age']] # Name: age, dtype: bool,
      ↳ ['city', 'gender', 'age'])' is an invalid key
```

```
[5]:   city gender age
      1   Delhi      F   22
      2  Mumbai      F   13
      3   Delhi      M   21
      4  Mumbai      M   12
      6  Mumbai      F   17
```

Update the values of a particular column on selected rows We can do this by running a for loop as well but if our dataset is big in size, then it would take forever to complete the task. Using loc in Pandas, we can do this within seconds, even on bigger datasets!

We just need to specify the condition followed by the target column and then assign the value with which we want to update

```
[6]: data
```

```
[6]:   age section   city gender favourite_color
      0    10      A  Gurgaon      M           red
      1    22      B   Delhi      F           blue
      2    13      C  Mumbai      F          yellow
      3    21      B   Delhi      M           blue
      4    12      B  Mumbai      M          black
      5    11      A   Delhi      M          green
      6    17      A  Mumbai      F           red
```

```
[7]: data.loc[(data.age >= 12),['section','age']]
```

```
[7]:   section  age
1         B   22
2         C   13
3         B   21
4         B   12
6         A   17
```

```
[8]: data.loc[(data.age >= 12),['section']] = 'M'
display (data)
```

	age	section	city	gender	favourite_color
0	10	A	Gurgaon	M	red
1	22	M	Delhi	F	blue
2	13	M	Mumbai	F	yellow
3	21	M	Delhi	M	blue
4	12	M	Mumbai	M	black
5	11	A	Delhi	M	green
6	17	M	Mumbai	F	red

1.1.4 Update the values of multiple columns on selected rows

```
[9]: # update multiple columns with condition
display (data.loc[(data.age >= 20), ['section', 'city', 'age']])
data.loc[(data.age >= 20), ['section', 'city']] = ['S', 'Pune']
display (data.loc[(data.age >= 20), ['section', 'city', 'age']])
```

	section	city	age
1	M	Delhi	22
3	M	Delhi	21

	section	city	age
1	S	Pune	22
3	S	Pune	21

1.1.5 Select rows with indices using iloc

```
[10]: data
```

```
[10]:   age section    city gender favourite_color
0    10         A  Gurgaon      M           red
1    22         S    Pune      F           blue
2    13         M  Mumbai      F          yellow
3    21         S    Pune      M           blue
```

4	12	M	Mumbai	M	black
5	11	A	Delhi	M	green
6	17	M	Mumbai	F	red

```
[11]: # select rows with indexes
data.iloc[[0,2]]
# 1st the row and the 3rd row
```

```
[11]:   age section   city gender favourite_color
0    10      A  Gurgaon      M              red
2    13      M   Mumbai      F             yellow
```

1.1.6 Select rows with particular indices and particular columns == Slicing

```
[12]: # select rows with particular indexes and particular columns
data.iloc[[0,2],[1,3]]
```

```
[12]:   section gender
0      A      M
2      M      F
```

Select a range of rows and columns using iloc

```
[13]: display(data)
data.iloc[1:3, 2:4]
```

	age	section	city	gender	favourite_color
0	10	A	Gurgaon	M	red
1	22	S	Pune	F	blue
2	13	M	Mumbai	F	yellow
3	21	S	Pune	M	blue
4	12	M	Mumbai	M	black
5	11	A	Delhi	M	green
6	17	M	Mumbai	F	red

```
[13]:   city gender
1    Pune      F
2  Mumbai      F
```

1.1.7 Operation 1: Use of Relational, logical and comparison operations

```
[14]: datasetExample = pd.read_csv('FinalOutput.csv')
display(datasetExample)
```

	eid	ename	esal	yearlySalary	department	UpdatedYearlySalary
0	1	Prashant	1000	12000	HR	13200.0

1	2	Amar	2000	24000	Ops	25200.0
2	3	Chitra	2000	24000	Admin	25200.0
3	3	Chitra	2000	24000	Admin	25200.0
4	4	Utkarsha	9878	118536	Ops	124462.8
5	5	Ajit	9999	119988	HR	125987.4

```
[15]: # 1. Extract the records of employees whose esal is greater than 5000
datasetExample.loc[datasetExample.esal > 5000]
```

```
[15]:      eid      ename  esal  yearlySalary department  UpdatedYearlySalary
4      4  Utkarsha  9878      118536      Ops      124462.8
5      5      Ajit  9999      119988      HR      125987.4
```

```
[17]: # 2. Extract only ename and department of employees whose esal is greater than
↳5000
datasetExample.loc[datasetExample.esal > 5000, ['ename', 'department']]
```

```
[17]:      ename department
4  Utkarsha      Ops
5      Ajit      HR
```

```
[18]: # 3. Extract those records whose sal is greater than 6500 and belongs to Ops
↳dept
# datasetExample.esal > 6500 # sal is greater than 6500
# datasetExample.department == 'Ops' # belongs to Ops dept

# (datasetExample.esal > 6500) & (datasetExample.department == 'Ops') # sal is
↳greater than 6500 and belongs to Ops dept

datasetExample.loc[(datasetExample.esal > 6500) & (datasetExample.department ==
↳'Ops')]
```

```
[18]:      eid      ename  esal  yearlySalary department  UpdatedYearlySalary
4      4  Utkarsha  9878      118536      Ops      124462.8
```

```
[22]: # 4. Display the name of employees whose salary is greater than 5000 and
↳belongs to HR dept
datasetExample.loc[(datasetExample.esal > 5000) & (datasetExample.department ==
↳'HR'), ['ename']]#, 'esal', 'department']]
```

```
[22]:      ename
5  Ajit
```

```
[28]: datasetExample.loc[(datasetExample.ename == 'Prashant') | (datasetExample.ename
↳== 'Chitra')] # or /
```

```
[28]:
```

	eid	ename	esal	yearlySalary	department	UpdatedYearlySalary
0	1	Prashant	1000	12000	HR	13200.0
2	3	Chitra	2000	24000	Admin	25200.0
3	3	Chitra	2000	24000	Admin	25200.0

2 Assignment

- Replace the salary of the employee as 7000 if the current salary is less than 1500

2.0.1 Operation 2 - Dealing with Duplicate Rows

```
[29]: display (datasetExample)
```

	eid	ename	esal	yearlySalary	department	UpdatedYearlySalary
0	1	Prashant	1000	12000	HR	13200.0
1	2	Amar	2000	24000	Ops	25200.0
2	3	Chitra	2000	24000	Admin	25200.0
3	3	Chitra	2000	24000	Admin	25200.0
4	4	Utkarsha	9878	118536	Ops	124462.8
5	5	Ajit	9999	119988	HR	125987.4

```
[30]: datasetExample.drop_duplicates(inplace = True)
display (datasetExample)
```

```
# Designed only for duplicate rows
# This will update the Dataframe
```

	eid	ename	esal	yearlySalary	department	UpdatedYearlySalary
0	1	Prashant	1000	12000	HR	13200.0
1	2	Amar	2000	24000	Ops	25200.0
2	3	Chitra	2000	24000	Admin	25200.0
4	4	Utkarsha	9878	118536	Ops	124462.8
5	5	Ajit	9999	119988	HR	125987.4

The major demerit of DropDuplicate is the index is never reset automatically which may impact the fetch cycle of the data when performing EDA or Statistical Modelling.

2.0.2 Operation 3 - Groupby in Pandas

```
[31]: data = {'Company': ['GOOG', 'GOOG', 'MSFT', 'MSFT', 'FB', 'FB'],
            'Person': ['Sam', 'Charlie', 'Amy', 'Vanessa', 'Carl', 'Sarah'],
            'Sales': [200, 120, 340, 124, 243, 350]}

df = pd.DataFrame(data)

display(df)
```

	Company	Person	Sales
0	GOOG	Sam	200
1	GOOG	Charlie	120
2	MSFT	Amy	340
3	MSFT	Vanessa	124
4	FB	Carl	243
5	FB	Sarah	350

```
[33]: # We want to look at average sales companywise
df.groupby('Company').mean()
```

```
[33]:      Sales
Company
FB      296.5
GOOG    160.0
MSFT    232.0
```

```
[35]: # We want to look at std sales companywise
df.groupby('Company').std()
```

```
[35]:      Sales
Company
FB      75.660426
GOOG    56.568542
MSFT    152.735065
```

```
[36]: df.groupby('Company').count()
```

```
[36]:      Person  Sales
Company
FB          2      2
GOOG        2      2
MSFT        2      2
```

```
[37]: df.describe()
```

```
[37]:      Sales
count    6.000000
mean    229.500000
```

```

std      100.899455
min      120.000000
25%      143.000000
50%      221.500000
75%      315.750000
max      350.000000

```

```
[38]: by_comp = df.groupby('Company')
```

```
[39]: by_comp.describe()
```

```
[39]:
```

	Sales							
	count	mean	std	min	25%	50%	75%	max
Company								
FB	2.0	296.5	75.660426	243.0	269.75	296.5	323.25	350.0
GOOG	2.0	160.0	56.568542	120.0	140.00	160.0	180.00	200.0
MSFT	2.0	232.0	152.735065	124.0	178.00	232.0	286.00	340.0

```
[40]: by_comp.describe().transpose()
```

```
[40]:
```

Company		FB	GOOG	MSFT
Sales count		2.000000	2.000000	2.000000
mean		296.500000	160.000000	232.000000
std		75.660426	56.568542	152.735065
min		243.000000	120.000000	124.000000
25%		269.750000	140.000000	178.000000
50%		296.500000	160.000000	232.000000
75%		323.250000	180.000000	286.000000
max		350.000000	200.000000	340.000000

```
[41]: by_comp.describe().transpose()['GOOG']
```

```
[41]:
```

Sales count	2.000000
mean	160.000000
std	56.568542
min	120.000000
25%	140.000000
50%	160.000000
75%	180.000000
max	200.000000

Name: GOOG, dtype: float64

2.0.3 Operation 4: Performing Merge Operations in Pandas

Merging is the Pandas operation that performs database joins on objects


```
[42]: dfExample1 = pd.DataFrame([[4, 'QA'], [1, 'HR'], [3, 'Dev'], [2, 'Ops']] ,  
    ↪columns=['eid', 'dept'])  
display (dfExample1)  
print ()  
dfExample2 = pd.DataFrame([[1, 'Prashant'], [2, 'Gokul'], [3, 'Guna']] ,  
    ↪columns=['eid', 'ename'])  
display (dfExample2)
```

```
    eid dept  
0     4   QA  
1     1   HR  
2     3   Dev  
3     2   Ops
```

```
    eid  ename  
0     1 Prashant  
1     2   Gokul  
2     3     Guna
```

```
[43]: resultDF = pd.merge(dfExample1, dfExample2)  
display (resultDF)  
  
# The merge worked in this case because both dataframe have the common column  
↪eid  
# that too with the same name
```

```
    eid dept  ename  
0     1   HR Prashant  
1     3   Dev   Guna  
2     2   Ops   Gokul
```

```
[44]: resultDF = pd.merge(dfExample2, dfExample1)  
display (resultDF)
```

```
    eid  ename dept  
0     1 Prashant  HR  
1     2   Gokul  Ops  
2     3     Guna  Dev
```

```
[45]: dept = pd.DataFrame([[4, 'QA'], [1, 'HR'], [3, 'Dev'], [2, 'Ops']] ,  
    ↪columns=['eid', 'dept'])  
emp = pd.DataFrame([[1, 'Prashant'], [2, 'Gokul'], [3, 'Guna']] ,  
    ↪columns=['empid', 'ename'])
```

```
display(dept)
print ()
display(emp)
```

```
   eid dept
0    4   QA
1    1   HR
2    3  Dev
3    2  Ops
```

```
   empid  ename
0      1 Prashant
1      2   Gokul
2      3    Guna
```

```
[46]: resultDF2 = pd.merge(dept, emp)
display (resultDF2)

# MergeError: No common columns to perform merge on.
```

```

↳
↳ -----
↳
↳ MergeError                                Traceback (most recent call↳
↳ last)
↳
↳ <ipython-input-46-d0f7545afc11> in <module>
↳ ----> 1 resultDF2 = pd.merge(dept, emp)
↳       2 display (resultDF2)
↳
↳ /usr/local/lib/python3.7/site-packages/pandas/core/reshape/merge.py in↳
↳ merge(left, right, how, on, left_on, right_on, left_index, right_index, sort,↳
↳ suffixes, copy, indicator, validate)
↳     84         copy=copy,
↳     85         indicator=indicator,
↳ ---> 86         validate=validate,
↳     87     )
↳     88     return op.get_result()
```

```

/usr/local/lib/python3.7/site-packages/pandas/core/reshape/merge.py in
↳ __init__(self, left, right, how, on, left_on, right_on, axis, left_index,
↳ right_index, sort, suffixes, copy, indicator, validate)
    618         warnings.warn(msg, UserWarning)
    619
--> 620         self._validate_specification()
    621
    622         # note this function has side effects

```

```

/usr/local/lib/python3.7/site-packages/pandas/core/reshape/merge.py in
↳ _validate_specification(self)
    1196             ron=self.right_on,
    1197             lidx=self.left_index,
-> 1198             ridx=self.right_index,
    1199         )
    1200     )

```

```

MergeError: No common columns to perform merge on. Merge options:
↳ left_on=None, right_on=None, left_index=False, right_index=False

```

```

[48]: resultDF2 = pd.merge(dept, emp, left_on = 'eid', right_on = 'empid')
display(resultDF2)

```

	eid	dept	empid	ename
0	1	HR	1	Prashant
1	3	Dev	3	Guna
2	2	Ops	2	Gokul

Left Join : Returns all rows from the left table, even if there are no matches in the right table

```

[51]: world_champions={'Team':['India','Australia','West Indies','Pakistan','Sri
↳ Lanka'], 'ICC_rank':[2,3,7,8,4],
        'World_champions_Year':[2011,2015,1979,1992,1996], 'Points':
↳ [874,787,753,673,855]}

chokers={'Team':['South Africa','New Zealand','Pakistan'], 'ICC_rank':
↳ [1,5,9], 'Points':[895,764,656]}

df1=pd.DataFrame(world_champions)
df2=pd.DataFrame(chokers)

display(df1)
print ()

```

```
display(df2)
```

	Team	ICC_rank	World_champions_Year	Points
0	India	2	2011	874
1	Australia	3	2015	787
2	West Indies	7	1979	753
3	Pakistan	8	1992	673
4	Sri Lanka	4	1996	855

	Team	ICC_rank	Points
0	South Africa	1	895
1	New Zealand	5	764
2	Pakistan	9	656

```
[52]: display (pd.merge(df1, df2, on = 'Team', how = "left"))
```

	Team	ICC_rank_x	World_champions_Year	Points_x	ICC_rank_y	\
0	India	2	2011	874	NaN	
1	Australia	3	2015	787	NaN	
2	West Indies	7	1979	753	NaN	
3	Pakistan	8	1992	673	9.0	
4	Sri Lanka	4	1996	855	NaN	

	Points_y
0	NaN
1	NaN
2	NaN
3	656.0
4	NaN

Right Join : Preserves the unmatched rows from the second (right) table, joining them with a NULL in the shape of the first (left) table

```
[55]: world_champions={'Team':['India','Australia','West Indies','Pakistan','Sri_
↳ Lanka'], 'ICC_rank':[2,3,7,8,4],
      'World_champions_Year':[2011,2015,1979,1992,1996], 'Points':
↳ [874,787,753,673,855]}
chokers={'Team':['South Africa','New Zealand','India'],'ICC_rank':
↳ [1,5,9], 'Points':[895,764,656]}

df1=pd.DataFrame(world_champions)
df2=pd.DataFrame(chokers)

display(df1)
```

```
print ()
display(df2)
print ()

display(pd.merge(df1, df2, on='Team', how='right'))
```

	Team	ICC_rank	World_champions_Year	Points
0	India	2	2011	874
1	Australia	3	2015	787
2	West Indies	7	1979	753
3	Pakistan	8	1992	673
4	Sri Lanka	4	1996	855

	Team	ICC_rank	Points
0	South Africa	1	895
1	New Zealand	5	764
2	India	9	656

	Team	ICC_rank_x	World_champions_Year	Points_x	ICC_rank_y \
0	India	2.0	2011.0	874.0	9
1	South Africa	NaN	NaN	NaN	1
2	New Zealand	NaN	NaN	NaN	5

	Points_y
0	656
1	895
2	764

Full Outer Join : Returns all records when there is a match in either left (table1) or right (table2) table records

```
[57]: world_champions={'Team':['India','Australia','West Indies','Pakistan','Sri_
↳ Lanka'], 'ICC_rank':[2,3,7,8,4],
      'World_champions_Year':[2011,2015,1979,1992,1996], 'Points':
↳ [874,787,753,673,855]}
chokers={'Team':['South Africa','New Zealand','Zimbabwe'],'ICC_rank':
↳ [1,5,9], 'Points':[895,764,656]}

df1=pd.DataFrame(world_champions)
df2=pd.DataFrame(chokers)

display(df1)
```

```
print ()
display(df2)
print ()

display(pd.merge(df1,df2, on='Team', how = 'outer'))
```

	Team	ICC_rank	World_champions_Year	Points
0	India	2	2011	874
1	Australia	3	2015	787
2	West Indies	7	1979	753
3	Pakistan	8	1992	673
4	Sri Lanka	4	1996	855

	Team	ICC_rank	Points
0	South Africa	1	895
1	New Zealand	5	764
2	Zimbabwe	9	656

	Team	ICC_rank_x	World_champions_Year	Points_x	ICC_rank_y	\
0	India	2.0	2011.0	874.0	NaN	
1	Australia	3.0	2015.0	787.0	NaN	
2	West Indies	7.0	1979.0	753.0	NaN	
3	Pakistan	8.0	1992.0	673.0	NaN	
4	Sri Lanka	4.0	1996.0	855.0	NaN	
5	South Africa	NaN	NaN	NaN	1.0	
6	New Zealand	NaN	NaN	NaN	5.0	
7	Zimbabwe	NaN	NaN	NaN	9.0	

	Points_y
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
5	895.0
6	764.0
7	656.0

Inner Join : Selects all rows from both participating tables if there is a match between the columns

```
[59]: world_champions={'Team':['India','Australia','West Indies','Pakistan','Sri_
→Lanka'], 'ICC_rank':[2,3,7,8,4],
```

```

        'World_champions_Year':[2011,2015,1979,1992,1996], 'Points':
        ↳[874,787,753,673,855]}
chokers={'Team':['South Africa','New Zealand','India'],'ICC_rank':
        ↳[1,5,9],'Points':[895,764,656]}

df1=pd.DataFrame(world_champions)
df2=pd.DataFrame(chokers)

display(df1)
print ()
display(df2)
print ()
display(pd.merge(df1, df2, on='Team', how='inner'))

```

	Team	ICC_rank	World_champions_Year	Points
0	India	2	2011	874
1	Australia	3	2015	787
2	West Indies	7	1979	753
3	Pakistan	8	1992	673
4	Sri Lanka	4	1996	855

	Team	ICC_rank	Points
0	South Africa	1	895
1	New Zealand	5	764
2	India	9	656

	Team	ICC_rank_x	World_champions_Year	Points_x	ICC_rank_y	Points_y
0	India	2	2011	874	9	656

2.0.4 Operation 5: Concat Operation

```

[60]: empExample = pd.DataFrame([[1,'Prashant'],[2,'Gokul'],[3,'Guna']] ,
        ↳columns=['empid','ename'])

empExample2 = pd.DataFrame([[4,'Nik'],[5,'Ashish'],[6,'Asha']] ,
        ↳columns=['empid','ename'])

display (empExample)
display (empExample2)

```

	empid	ename
0	1	Prashant

1	2	Gokul
2	3	Guna

	empid	ename
0	4	Nik
1	5	Ashish
2	6	Asha

```
[61]: # Rowwise Concatenation --- Ensure the column names are same in all DFs.
```

```
resultEmp = pd.concat( [empExample,empExample2])
resultEmp
```

```
[61]:
```

	empid	ename
0	1	Prashant
1	2	Gokul
2	3	Guna
0	4	Nik
1	5	Ashish
2	6	Asha

```
[62]: resultEmp = pd.concat( [empExample,empExample2] , axis = 1)
resultEmp
```

```
[62]:
```

	empid	ename	empid	ename
0	1	Prashant	4	Nik
1	2	Gokul	5	Ashish
2	3	Guna	6	Asha

```
[63]: empExample3 = pd.DataFrame([[4,'Nik'],[5,'Ashish'],[6,'Asha']] ,
↳columns=['empid','empname'])
display (empExample3)
```

	empid	empname
0	4	Nik
1	5	Ashish
2	6	Asha

```
[64]: resultEmp2 = pd.concat( [empExample,empExample3] , axis = 0)
resultEmp2
```

```
[64]:
```

	empid	ename	empname
0	1	Prashant	NaN
1	2	Gokul	NaN
2	3	Guna	NaN
0	4	NaN	Nik

1	5	NaN	Ashish
2	6	NaN	Asha

```
[65]: empExample3.columns = ['empid','ename']
display (empExample3)
resultEmp2 = pd.concat( [empExample,empExample3] , axis = 0)
resultEmp2
```

	empid	ename
0	4	Nik
1	5	Ashish
2	6	Asha

```
[65]: empid      ename
0      1  Prashant
1      2    Gokul
2      3     Guna
0      4      Nik
1      5   Ashish
2      6    Asha
```

```
[ ]:
```